

The Value Realization Score

One North Star for product adoption and customer value capture

Erin Kim · Principal PM, Product Analytics

Today We Measure Consumption — Not Value

Today, product adoption is measured on pure monthly consumption — a single ratio that ignores churn-risk behaviors and routinely misses expansion opportunities. The goal is a framework that shows whether customers are truly getting value from what they purchased, beyond monthly consumption — which means detecting Spike & Drop, Shelfware, and Consistent Overage, and adjusting correctly for mid-year expansions.

Spike & Drop

~5% of accounts

Burn 90% of the entitlement in Q1, then flatline. Cumulative utilization calls them a star account — they are churning.

Shelfware

~10% of accounts

Bought, never turned on. Consumption-only dashboards literally cannot see them — zero rows, zero alarms.

Consistent Overage

~15% of accounts

Run at 120%+ of license. A linear metric ranks them healthiest of all — and hides the uncaptured upsell.

Mid-Year Expansion

6 accounts

A second, larger contract mid-year, creating overlapping active entitlements. Naive math double-counts the license and flags a healthy, growing account as shelfware.

Measure Value Through Value Realization

$$\text{VRS} = 100 \times (0.35 \cdot \text{Utilization} + 0.25 \cdot \text{Feature Adoption} + 0.25 \cdot \text{Sustained Usage} + 0.15 \cdot \text{Time to Value})$$

35% License Utilization

Depth — consumed ÷ purchased (every active contract summed, so mid-year expansions count correctly), normalized 0–100%. Near-zero consumption exposes shelfware; consumption above the license raises the consistent-overage flag.

25% Feature Adoption

Breadth — average depth across every entitled feature, normalized 0–100%. A flag catches single-feature dependency inside high-utilization accounts.

25% Sustained Usage


Durability — recency-weighted usage over the trailing 3 months; detects lack of sustained usage (spike & drop).

15% Time to Value

Speed — time to first usage, in ≤30 / 30–90 / 90+ days.

Scores combine by revenue, not by count — a customer's, platform's, or portfolio's score is the average of its product scores, weighted by each product's ARR. A struggling \$900K product moves the score 18× more than a healthy \$50K add-on.

From Raw Events to One Number, Reproducibly



● Data Quality

The dataset mirrors the real catalog — 19 PANW product models, feature volumes at realistic scale — with the four anomaly behaviors injected at known rates as hidden ground truth. Quality is enforced, not assumed: 36 automated checks verify table schemas, every contract has usage rows, monthly license totals equal the sum of active contracts, and the injected rates are recovered exactly.

● Design

Each signal was built to catch a specific problem: License Utilization catches shelfware and overage, Sustained Usage catches spike & drop, Feature Adoption catches accounts riding on one feature, Time to Value flags slow starts — and mid-year expansions are handled in the contract math. Automated checks confirm every planted problem is caught, with zero false alarms, and every catch is measured in dollars.

● Model

Simple, readable scoring rules instead of machine learning: a usage curve that rewards real use and caps over-use, feature scoring against each feature's normal volume, extra weight on recent months, and a fixed order of rules that gives each account one status. Chosen because this score will drive pay — everyone must see exactly why a number is what it is. ML comes later, to tune the weights once real renewal outcomes exist.

The Portfolio Today — June 2026

\$146.5M

Total ARR

70.2

Portfolio VRS (ARR-wtd)


\$28.4M

ARR at risk — 19.4%

\$20.3M

Expansion ARR

ARR at Risk by Platform (\$M)



- **SASE — \$14.3M: driven by shelfware and spike & drop**
\$7.6M in shelfware (seats never rolled out) and \$6.7M in spike & drop (used early, then went quiet). Nearly all of it is Prisma Access Mobile Users (\$5.9M) and Remote Networks (\$5.4M) — likely seats bought ahead of migrating users.
- **Hardware NGFW — \$12.8M: mostly shelfware, plus spike & drop**
\$7.9M in shelfware (licenses never activated) and \$4.9M in spike & drop (usage stopped completely). Over half sits in just three models — PA-5410, PA-5445, PA-410 — a short, targeted rescue list.
- **Cloud NGFW — \$0.7M: mostly shelfware, and Azure lags AWS**
\$0.5M shelfware, \$0.2M spike & drop — 19% of its ARR at risk, a bigger share than hardware or software. Azure customers average a score of 60 vs 76 on AWS: fix the Azure path while cloud is small.
- **Software NGFW — \$0.6M: the smallest risk**
Just 14% of its ARR at risk — the lowest of any platform — from a little shelfware and one spike & drop account. These customers adopt fastest and keep using what they buy.

Opportunities to Address ARR Risk and Drive Expansion

ACTIVATE

\$16.4M

Shelfware Risk — bought, never deployed. Customer Success runs first-deployment plays; unrescued accounts are flagged renewal risk early.

Customer Success

WIN BACK

\$12.0M

Spike & drop — burned early, went dark. Diagnose why usage stopped and build a save plan while the renewal is still winnable.

CS / Account team

UPSELL

\$20.3M

Expansion flag — consuming past the license. Real demand, currently uncaptured revenue. Hand the list to GTM for capacity true-ups.

GTM

PRODUCT — invest in driving feature adoption. It is the weakest signal on all 19 product models — customers use about half the features they own. It rarely puts an account at risk by itself; it is what caps healthy accounts near 70 instead of 90 — the biggest headroom on the North Star. Start with guided onboarding, feature-discovery nudges, and “light it up” campaigns for entitled-but-idle features.

CUSTOMER SUCCESS & GTM — measure GTM teams on VRS. Pay on % of ARR in Value Realized and net state improvements — never raw consumption. An overage account credits both teams correctly: CS for healthy adoption (score capped), GTM for the expansion flag.

Decisions We Need From This Room

1

Backtest before full rollout

Validate VRS against historical data to confirm that low scores correlate with churned customers, high scores correlate with renewals/expansions.

2

Adopt VRS as the North Star

One holistic score to measure if customers are getting value from what they purchased, beyond just consumption.

3

Tie incentives to VRS

Measure CS/GTM teams on VRS and compensate them on it.

Live demo: the VRS dashboard — then discussion. Every threshold shown today is a parameter, and every number is reproducible.

Appendix — How the Pipeline Is Tested

Category	What was tested	Checks
Unit	Scoring logic against planted ground truth: every injected anomaly lands in its intended state or flag, zero false positives on clean accounts — plus curve properties (shelfware scores 0, overage capped at 0.8 and flagged, partial time-to-value only in the 31-90-day band).	13
Data Quality	Table health: no orphaned records across any join, unique customer × product × month grain, values in range, no zero licenses, all 12 months present, every contract has usage rows, latest month is current.	11
Schema	Expected columns and types asserted on all 6 raw tables and the scores mart — guards against silent load and auto-detect drift (a failure mode hit once during a BigQuery reload, now caught by name).	7
Data Integrity	Source→target: monthly licensed totals independently recomputed from raw contracts must match the pipeline with zero mismatches — proving mid-year expansion resolution.	2
Regression	Aggregate checksums pinned: total row count, portfolio VRS checksum, latest-month ARR. Any unintended change to existing data fails the run.	3
Cross-Env	Parity between live BigQuery and the offline CSV build: row count, VRS and ARR checksums compared on every run (non-fatal, advisory).	3

Coverage: every required anomaly maps to at least one test. 36 fatal assertions + 3 parity checks run as a single gate against live BigQuery — currently 36/36 green.